

Chelsea Dole

+1 (206) 914-7021 ▪ chelseadole@gmail.com

[Github](#) ▪ [LinkedIn](#) ▪ [Website](#)

TECHNICAL SKILLS

- | | | |
|--|---|---|
| <ul style="list-style-type: none">• Python 3• JavaScript• Node.js• Django• Flask | <ul style="list-style-type: none">• PostgreSQL• Elasticsearch• Apache Kafka• RabbitMQ• Kubernetes | <ul style="list-style-type: none">• Cloud computing (GCP, AWS)• Docker• Linux• Git |
|--|---|---|

WORK EXPERIENCE

Coffee Meets Bagel, *Remote*

Data Engineer III

April 2020–Present

- Migrated a self-managed Elasticsearch cluster to AWS Elasticsearch Service. Refactored data model resulting in 15% monthly cost savings, 11x reduction in total index document count, and 35% reduction in max query time
- Organized and led several organization-wide projects aimed at improving code quality and system transparency, including: an automated release system for internal Python packages, the transition from schema-less to Avro internal events, and the creation of backend system runbooks
- Currently designing and implementing a company-wide migration of all ETLs and events from RabbitMQ to Apache Kafka

Data Engineer II

June 2019–April 2020

- Designed and fully implemented an event-based feature engineering system providing real-time features to CMB's internal user-matching pipeline. Resulted in 85% reduction in feature calculation time, estimated 75%+ reduction in engineering time needed to deploy/backfill new features, and vastly increased visibility into model inputs
- Supported data scientists by building backend architecture to train and deploy user matching models

Technologies: Python, PostgreSQL, Elasticsearch, Kafka, AWS (EC2, ECS, RDS, MSK, Redshift, Kinesis, S3, etc), Django, RabbitMQ, Kubernetes, Redis/Memcached, Jenkins, Docker, DataDog, Git

Zonar Systems, *Seattle*

Software Development Engineer I

Feb 2018–June 2019

- Built and scaled backend applications to support a data analysis web platform for Fleet Pipeline team.
- Developed components of a Python/Flask data analysis API, a Golang event logging API, and a Dataflow data summarization pipeline

Technologies: Python, Golang, JavaScript/Node.js, Kubernetes, GCP (Dataflow, BigQuery, etc), Flask, Pandas/NumPy, Angular 2, PostgreSQL, Redis, DataDog, Git

PROJECTS // SPEAKING

[Understanding Python's Memory Management](#)

Mar 2021

10-minute introduction to CPython's internal memory manager given remotely as part of PyLadies International Women's Day 2021 event

[Dreaming of Databases: 5 Questions for Data Layer Decisions](#)

Jan 2020

25-minute talk given at the ThaiPy meetup in Bangkok to 100+ technical attendees. Talk covered database vocabulary, and discussed how to make educated decisions regarding database choice

[Going Serverless: Deploying with GCP App Engine and Cloud Functions](#)

Jan 2019

30-minute talk given at Puget Sound Programming Python (PuPPy) meetup to 120+ software engineers. Talk included the live deployment of a demo API with App Engine and the gcloud SDK

[Personal Portfolio](#)

Jan 2018

Simple personal website using HTML, CSS, and Javascript. Hosted using Github Pages

[write-me](#)

Dec 2017

PyPI package that generates elegant, specific READMEs based on a user's Python codebase.

Technologies: Python, Sphinx, PyPI

EDUCATION // CERTIFICATIONS

Code Fellows | Seattle, WA

Aug–Dec 2017

Certificate — Advanced Software Development in Python

University of California, Berkeley

2013–2017

Degree: Political Economy (B.A.)

Coursework: Advanced economics, statistics, computer science, political history/theory, business, int'l trade

GPA: 3.40

Other:

- *Learn Apache Kafka* (Udemy, 7 hours)
- *Kafka Cluster Setup & Administration* (Udemy, 4 hours)